

Sanjeev Kumar

Advanced Big Data Analytics for Business with R

dedicated to the one to whom I owe all

Contents

List of Tables	v
List of Figures	vii
Go Blue!	ix
Preface	xi
Part II: Data Exploration and Visualization	xiv
Part III: Traditional Statistical Modeling	xiv
Part IV: Machine Learning and Predictive Analytics	xv
Part VI: Appendices, Bibliography and Index	xv
About the Author	xix
1 Introduction	1
1.1 Archive - Delete	1
2 Advanced Big Data Analytics	3
3 Our Tools: R, RStudio	5
3.1 Introduction to R and RStudio	5
3.1.1 Download and Install R and RStudio	5
3.2 Useful Features of RStudio	7
3.2.1 Console	7
3.2.2 Editor	8
3.2.3 Environment	8
3.2.4 History	9
3.2.5 Files	9
3.2.6 Plots	9
3.2.7 Packages	10
3.2.8 Help	10
3.3 Helpful Menu Items and Keyboard Shortcuts	10
3.4 Essential R Tips and Tricks	10
3.4.1 R is Case Specific and Space Neutral	10
3.4.2 Comments and Line Continuation	10
3.4.3 Naming Conventions	11
3.5 RMarkdown (RMD) Files	11

3.6	Embedding R Code Blocks in RMD Files	12
3.6.1	Embedding Options	12
3.6.2	Embedding Code Blocks of Other Languages	13
3.6.3	Embedding in-line R Code	13
3.6.4	Reproducibility of Data Analysis	13
3.7	Key Features of R/RStudio Ecosystem	13
3.7.1	Open Source Software	13
3.7.2	Extensability of R	14
3.7.3	Usability of R	14
3.7.4	R is NOT Verbose	15
3.7.5	Object Oriented Programming	15
3.7.6	R Community's Sense of Humor	15
3.8	Common Issues, Pitfalls, Bugs, Obstacles and Their Solutions	16
3.8.1	I have an older R installtion. How do I upgrade it? . .	16
3.8.2	Help - I am not being able to install any packages! . .	16
3.8.3	Why is the edit command not working for me?	17
Appendix		19
A Syllabi		19
A.1	TO404 Big Data Manipulation and Visualization	19
A.2	TO414 Advanced Analytics	19
A.3	TO628 Advanced Big Data Analytics	19
Bibliography		21

List of Tables

1.1 The boring iris data.	1
-----------------------------------	---



List of Figures

1	ix
1.1	Hello World!	2
3.1	Staring Window of Base R Installation	6
3.2	Staring Window of RStudio	7
3.3	RStudio Editor Window	8
3.4	Learning Curves for R (red) and alternatives (blue)	15



Go Blue!



FIGURE 1

Until I build a cover image for this book, Woodson's Giant Leap will have to suffice.



Preface

We live in a world awash in data. Companies are increasingly turning to data analytics to extract a competitive edge from data, especially large, complex datasets often called *Big Data*. However, companies are increasingly facing the challenge posed by the scarcity of analytical talent - people who can turn data into better decisions, people who can extract insights and information from data. There is growing demand for professionals with strong quantitative skills combined with an understanding of how data analytic techniques can be applied to business contexts and managerial decision making. To help my students succeed in this growing field, I teach classes in **Advanced Big Data Analytics** in Ross School of Business, Univ. of Michigan. In these courses I teach advanced analytical, statistical and data mining tools with an applied focus. This book is developed specifically for these courses.

The main focus of this book (and the associated courses) is to prepare students to model and manage business decisions with data analytics and decision models using real life case contexts and datasets. By the end of this book students will have a better understanding of processes, methodologies and tools used to transform the large amount of business data available into useful information and support business decision making. The book will focus on extracting actionable business intelligence by analyzing traditional business data as well as more recently available datasets such as social media, crowdsourcing and recommendation engines. The book focuses on the powerful, open source (and hence free) data analysis environment R.

As I designed these courses, I realized that while there are a lot of references available for doing Advanced Analytics on Big Data using R, there isn't a good reference that approaches the topic from the perspective of business students and is accessible for students who do not have extensive background in statistics or computer programming. I designed my classes to have an applied nature with significant amount of hands-on work on real business Big Data with significant managerial implications. I emphasized aspects of the analytics process that are important for actual practice of data science but are not typically well covered in traditional textbooks - like data cleaning, managing large datasets and building data dashboards. I ended up creating a significant amount of material for the classes - much of it collated from already available but widely dispersed sources. This book is a collection of these material.

Business students have a unique mixture of tech savvy, super smarts and learn-

ing ability with a relative lack of computer programming or coding experience. They further have a great applied sense of statistics and data analysis but typically without the theoretical expertise of a Statistics student. This book is directed towards such students. This book assumes that business students are joining an Advanced Analytics course without any knowledge of computer programming and without any background in R. This book assumes that students are familiar with basic probability and statistics but their focus is on applied statistics - not on the theory. The book then builds up student's comfort level with R while at the same time making progress on key Advanced Analytics materials.

Direct all feedback on this book to the author at email: `sankum at umich.edu` or twitter: `@a_teachr`.

Who Is This Book For?

This book is for business students, practitioners and executives who want to have hands-on experience in working with business related big data and want to extract actionable insights from the data using advanced analytical techniques. This book is application oriented and hence focuses on the application of advanced analytical techniques rather than the theoretical details. Consequently, this book is not for readers solely focused on the theoretical, statistical part of data analytics.

This book assumes that the readers have basic familiarity with introductory probability and statistics. This domain is easier to follow and understand if the reader also has *some* exposure to *some* kind of computer programming although that is not a necessity. This book has been written for a practitioner audience, not an academic one. The book aims to be an exhaustive reference resource for practitioners - that's why it starts with baby steps of introducing R and basic data manipulation and ends at the other end with advanced Machine Learning algorithms and complex model improvement algorithms.

Thank you for placing your trust in this book. This chapter will provide further details about this book and how to best read/use this book.

About This Document

This document, which you are probably seeing as a GitHub eBook or a PDF document or even as printed book, has been written entirely using R and RStudio. This book is built using the **knitr** package (Xie, 2015) and the **bookdown** package (Xie, 2017). PDF version of the book is created using my favorite document creation system called **LaTeX**. This document has been created entirely using Open Source tools and has been released back into the Open Source ecosystem for free using the GNU General Public License (GPL). You are free to share this document with others as long as you comply with the GPL license. GPL License usually require that the product (in this case, this book) be made available free or charge; and any subsequent product made using the GPL Licensed product must also be made available free of charge (Foundation, 2007).

This is my first attempt at writing a book size document. I have had to develop a bunch of workarounds to make the process work and get a quality output. The source code for this book is available for public use free of charge at: <https://github.com/clarifyR/AABookDown>. The HTML and PDF versions are available (again, free of charge) at: <http://clarifyr.com/AABookDown/>. This book is a live document which is continually updated as I teach my classes and learn more about student needs. A printed version will be available for purchase from Amazon once the book is substantially complete.

Following is the R session information that built the current version:

```
xfun::session_info()

## R version 3.4.3 (2017-11-30)
## Platform: x86_64-apple-darwin15.6.0 (64-bit)
## Running under: macOS 10.14.1
##
## Locale: en_US.UTF-8 / en_US.UTF-8 / en_US.UTF-8 / C / en_US.UTF-8 / en_US.UTF-8
##
## Package version:
##   backports_1.1.2 base64enc_0.1.3 bookdown_0.5
##   compiler_3.4.3  digest_0.6.13  evaluate_0.10.1
##   graphics_3.4.3  grDevices_3.4.3 highr_0.6
##   htmltools_0.3.6 jsonlite_1.5   knitr_1.18
##   magrittr_1.5     markdown_0.8   methods_3.4.3
##   mime_0.5         Rcpp_0.12.14   rmarkdown_1.8
##   rprojroot_1.3-2 rstudioapi_0.7 stats_3.4.3
##   stringi_1.1.6   stringr_1.2.0  tools_3.4.3
```

```
##   utils_3.4.3      xfun_0.3      yaml_2.1.19
```

How is This Book Structured

This book covers a wide range of material. The material is organized in 5 parts, 19 chapters and several appendices.

Part I: Getting Started

We set ourselves for the book - download and install all needed software; figure our way around R and RStudio and learn the basics of the R language. Essentially lay down enough of a foundation that we can start getting productive. For students without a computer programming background, it is essential that they do not rush this part. Our success in later parts depend upon us getting comfortable with the material here.

Part II: Data Exploration and Visualization

Real data analytics begins with us understanding and getting a handle on the data. This typically involves cleaning up the data, generating descriptive statistics to better understand the data and finally creating visualizations that allow us to better understand the underlying complexity of the data. In my opinion, data cleaning is the most under-appreciated part of data analytics. It often takes more time and effort than the actual analysis that follows.

Data visualization has emerged as a key tool in Big Data Analytics. We specifically focus our attention here on a subset of data visualization that is important in the business context - creating data dashboards that can help managerial decision making.

Part III: Traditional Statistical Modeling

Even advanced analytic techniques like machine learning algorithms in the next part have their foundations in traditional statistical methodologies. Classical statistical modeling approaches like Linear Regression are still the benchmark given their immense popularity, flexibility and stability. In this part of the book, we explore traditional statistical analysis tools like Linear Regres-

sion, Generalized Linear Models like Logistic and Survival Models, Principal Components and Factor Analysis, Time Series Analysis and so on.

As we assume that you are already familiar with basic probability and statistics, we will focus on application of the methodologies and not the underlying theory. We will also focus on how to tweak these tools for the Big Data world as many of these tools run into trouble when sample sizes are quite large. Bigger is not always better - traditional statistical methodologies were developed/optimized for smaller sample sizes. Using them for large sample sizes give rise to unique issues and problems - we will discuss how to address them.

Part IV: Machine Learning and Predictive Analytics

This is the largest and the most important part of the book. This is why this book was written in the first place - an overview of data analytics tools specific to Big Data - variously known as Machine Learning, Statistical Learning, Predictive Analytics, Data Mining as so on. This book focuses on tools that help us make sense of Big Data and helps us automate the extraction of managerial decision making insights from Big Data.

As with much of the rest of the book, the focus is on applying the tools rather than their theoretical/mathematical underpinnings. We will discuss enough theory to develop an overall understanding and then devote our energies on making these tools work on real datasets.

Part V: Putting It All Together

Now that we have gone through all the elements individually, we can move forward to create a combined, integrated approach that puts all these pieces together. How to combine different models so that they result in an output better than sum of their parts? How to ensure that our models improve as more data become available?

Part VI: Appendices, Bibliography and Index

Back matter of the book. The Index in the end has three main components - Key Concepts, R Commands and R Packages. Appendices include syllabus for the courses this book is primarily used for and other miscellaneous content that did not fit in one of the main matter chapters. The index is only available in the PDF format of the book.

There are a lot of quality, free, online resources available for building up your R and Data Analytics skills before you go through this book or the associated

courses. They are listed in Appendix: Key References. A fully fleshed data analysis example has been presented on Appendix: Data Analysis Example to give you an idea of the power and range of what can be accomplished with just a little bit of familiarity with R.

How to Read This Book

You would have noticed by now that this book integrates R code within its text. Most of the time R code takes the form of dedicated code blocks like the one below:

```
#This is a demo code block  
print("Hello World")
```

```
## [1] "Hello World"
```

As you can see from above, code blocks are printed in color, in **fixed width font**. There is a color scheme here that you will soon become familiar with - comments, functions, strings all have their specific formatting. The output of the code block appears right after - for example - the output of the **print()** command follows the code above.

You would notice that whenever we refer to a R command in a significant way (like **print()** in paragraph above), the command is printed in **fixed width font**, in bold that makes it easy for you to see which commands are discussed significantly on the page. The highlighting also ensures that an entry for the command is placed in the Index provided at the end of the book. For times when a command is mentioned in text in a minor way not necessitating an Index entry, they will be typed in **fixed width font**.

Like R Commands, this book provides special formatting for R Packages used in the book. R Packages are formatted like R Commands with an entry in the Packages section of the Index. Lastly, special formatting is provided for key concepts - similar to R Commands and Packages and their Index entry is in the Key Concepts section. For example: this book focuses on **Open Source** software - that are created by a community of software developers for use by the community, usually made available for free.

The three special formatting elements - commands, packages and key concepts - ensure that you have a summary of significant elements discussed in the page just by looking for highlighted items.

Is This Book Suitable For You?

Well - you wouldn't know until you spend some time with it. Dig in.

A quick word of caution though before you get too deep: this book (and the associated courses) are very hands-on. There is no point in reading this book like a sequence of text. This book should be seen more as an illustrated text - illustrated with relevant R commands and material. You should read this book alongside an RStudio session - trying all the commands there as you read along. You will need to get comfortable with a lot of command line typing, keyboard shortcuts and figuring workarounds to inevitable problems that will arise. We will often get into situations where there will be no tested/optimized/prescribed solution and we will need to figure our way out - often with some trial and error. You should be comfortable with such ambiguity.

Acknowledgments

Add acknowledgements here.

Sanjeev Kumar
Technology and Operations, Ross School of Business, Univ. of Michigan



About the Author

Sanjeev Kumar is part of the Technology and Operations faculty at the Ross School of Business, Univ. of Michigan, Ann Arbor.



{-}



1

Introduction

Welcome to **Advanced Big Data Analytics for Business with R**. Let's get started.

1.1 Archive - Delete

This section holds the placeholder information that will be deleted after the book is built.

We have a nice figure in Figure 1.1, and also a table in Table 1.1.

```
par(mar = c(4, 4, 1, .1))
plot(cars, pch = 19)
```

```
knitr::kable(
  head(iris), caption = 'The boring iris data.',
  booktabs = TRUE
)
```

TABLE 1.1: The boring iris data.

Sepal.Length	Sepal.Width	Petal.Length	Petal.Width	Species
5.1	3.5	1.4	0.2	setosa
4.9	3.0	1.4	0.2	setosa
4.7	3.2	1.3	0.2	setosa
4.6	3.1	1.5	0.2	setosa
5.0	3.6	1.4	0.2	setosa
5.4	3.9	1.7	0.4	setosa




FIGURE 1.1: Hello World!

2

Advanced Big Data Analytics

We talk about the *FOO* method in this chapter.

‘In God We Trust; All Others Must Bring Data’ – *William Edwards Deming*



3

Our Tools: R, RStudio

R is awesome. R is also considered pretty difficult to learn. In my opinion, that's not true. In this chapter, we will quickly get our feet wet in the shallow end of the R pool - ready to jump off the deep end in the succeeding chapters.

3.1 Introduction to R and RStudio

For the purpose of this course, R is a Statistical Computing Environment. It provides us with an underlying computer language, analysis tools and software platform for us to work with data. While R is pretty powerful, it is built by geeks for geeks. As a result, it lacks the user-friendliness needed by beginners and casual users. To help mitigate this problem, we do not directly interface with R and instead use the user friendly RStudio software as a wrapper around R. RStudio is often considered an IDE (Integrated Development Environment) for R.

For the record, R is a pretty awful name for anything - try googling "R" and see how relevant the search results are! It is called R after the first name of the two authors of R: Robert and Ross. The name kinda makes sense when you consider that R is based on another programming language named S.

3.1.1 Download and Install R and RStudio

R can be downloaded from any of the mirror sites maintained by the Comprehensive R Archive Network ("CRAN"). I recommend the CRAN Mirror maintained by RStudio: <https://cran.rstudio.com/>. Here you will find links to download R for Windows, Mac-OS and Linux. The current version of R available for Windows is R 3.4.3 and measures at around 70MB. Follow the directions to install the starting packages for R - often called "Base R".

Even though we will not use a bare R installation - let's take a look at it anyway. Launch R once the installation finishes. It will look similar to the figure below.

This window is called the R Console. As you would note - it looks a little scary - no icons, nothing clickable, no menu even. While it is possible for us to do everything we wish to do with R just with the R Console, it really would be nice to make our experience with R a little more user friendly. Lucky for us, we have RStudio to pick up the slack here.

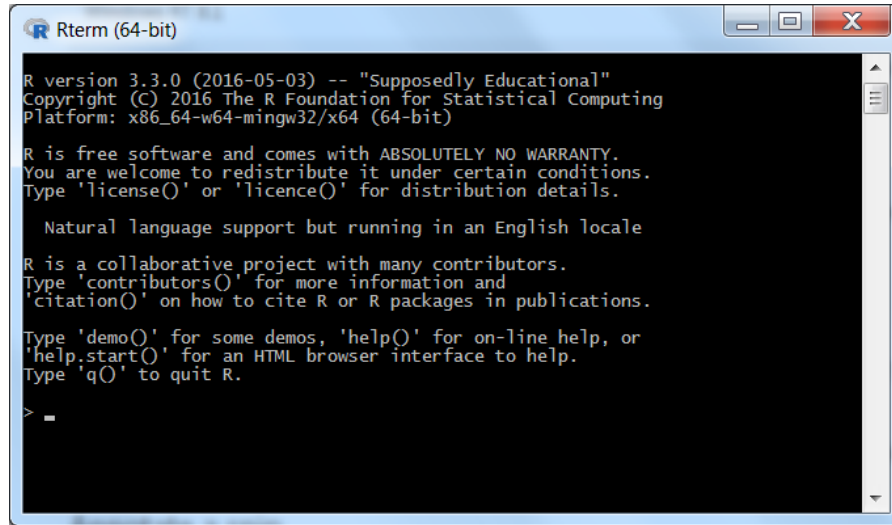


FIGURE 3.1: Staring Window of Base R Installation

Once you have finished installing R, you can download and install RStudio from: <https://www.rstudio.com/products/rstudio/download/>. There again are installers available for Windows, Mac-OS and Linux. After you are done installing RStudio, you should launch RStudio to check that everything is installed and connected properly. When you launch RStudio for the first time, you will see something similar to the figure below.

The larger part of the RStudio window is same as the R Console and this is in fact R running inside RStudio. When we launched R independently, this is the prompt that you we got. RStudio combines this direct access to R with other useful elements. In the bottom right, you have access to the file system, you can manage R packages (more about them later), read Help documentation and see various plots and charts that you will soon be building. In the top right, RStudio provides you with details of current R environment and a history of your recent R commands. If you are using Git repository for your files then you will see Git tab here.

Now that we have both R and RStudio set up, it is time for us to get started in making some use of it. Note that the R environment is used primarily by typing commands on the Console prompt or running commands saved in a file - typically a R Script file or a RMarkdown file. While RStudio provides GUI

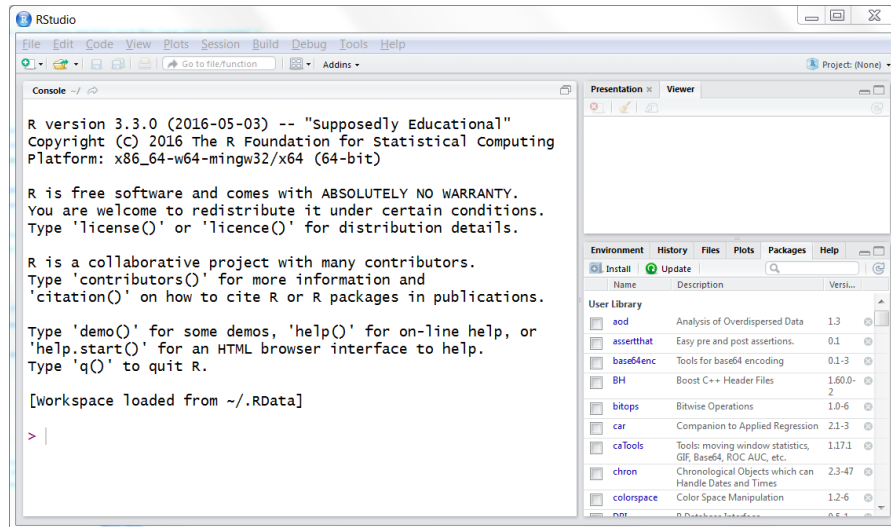


FIGURE 3.2: Staring Window of RStudio

elements for many common tasks, we will emphasize corresponding command line versions which usually provide better control and flexibility compared to their GUI counterparts.

3.2 Useful Features of RStudio

3.2.1 Console

The console is where you can interact directly with the R system. You can write one line at a time and immediately execute it there by pressing enter. If you wish to write more than one command in one line then you can separate the commands using the semi-colon character. The console is essentially an R session running inside RStudio. The console shows a command prompt (usually the `>`) character to indicate that you can type there. You can change the prompt by changing the options setting.¹

¹I once thought it was very clever to have this as my prompt “Sanjeev Says:”. I no longer think so.

3.2.2 Editor

The editor is where you can work on and edit all of your files. This is a pure text editor - so no formatting of any kind. You will mostly be working with RMD files here. The tab panel at the top of the editor allows you to work on multiple files at once. The editor window is hidden when you start RStudio as you usually don't have a file to edit when you begin. Once you create/open a new file then they will open in the editor window and you can edit them. The editor window helpfully shows you the line numbers in left column and the cursor positioning in the line at the bottom left corner. The bottom right corner shows you the file type that you are working on. There are usual icons for saving a file, finding, spell check etc on the top bar. See Figure: ?? for an example.

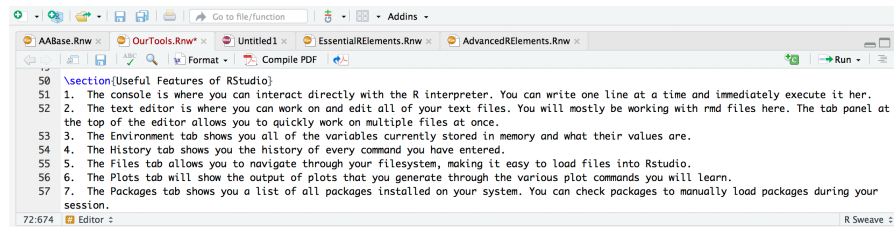


FIGURE 3.3: RStudio Editor Window

3.2.3 Environment

The Environment tab shows you all of the variables currently stored in memory and what their values are. We haven't said what variables are as yet. Hold on - its coming in the next chapter. The tab in fact shows all the objects that currently exist in your R session. Objects will be discussed very soon, later in this chapter. Environment tab is a very useful place to have a quick look at the state of your current R session. The environment tab has a handy broom icon for deleting all the objects that exist and clearing your environment. You can, of course, do that using a command as well - `rm` for remove or delete. The command `ls` generates a list of all the objects in the environment. We can combine the two to delete every object in the environment if we so choose:

```
rm(list = ls()) #Delete every object in the environment
```

3.2.4 History

The History tab shows you the history of all the commands you have entered in the Console window. You can use available icons on top to search through the history or delete items from the history. Two very useful icons are for transferring commands to the console and to the source. See details in figure below. . The **To Console** option allows you to select one or more commands from the history and transfers them to the Console window where you can run them by pressing Enter. The **To Source** option sends the selected commands to the Editor window, to the currently open and active file.

```
# %[History Window](images/history.png)
```

3.2.5 Files

The Files tab allows you to navigate through your filesystem, making it easy to load files into Rstudio. It shows the content of the current working directory when RStudio starts. You can navigate through the file system of your machine using the available options. The top icons provide ability to create a new folder, delete selected files, rename selected files and refresh file listings. The **More** option reveals further functionalities including two useful ones: the ability to set the currently displayed directory as the working directory and the ability to get the Files tab to display the contents of the current working directory. Figure: shows the options.

```
# %[Files Window](images/files.png)
```

3.2.6 Plots

the Plots tab shows the plots generated by the commands typed in the Console. See figure show the plot generated by the `plot(cars)` command in the Plots tab. As you can see, you can use options to zoom in/out and export the plot image as an image file or a PDF document. Plots made by R plots are Vector graphics and hence should be saved in vector formats like PNG and not bitmap formats like JPG.

```
# %[Plots Window with a Demo Plot](images/plots.png)
```

3.2.7 Packages

The Packages tab shows you a list of all the packages installed on your system and their version number. You can check the checkbox next to the package name to manually load packages during your R session. You will also find icons for installing and updating packages. A search icon allows for searching among installed packages.

3.2.8 Help

As the name suggests, the Help tab can be used to find Help documentation on specific keywords. You can search for a keyword in the search field for looking at corresponding help documentation. If you write a help command in the console, then the output of that will also show in the Help tab. For example, if you type the command below (looking for help on “help”), you will get the help documentation as shown in figure below.

```
help("help")  
# %![Help Window when looking for help on "help"](images/help.png)
```

3.3 Helpful Menu Items and Keyboard Shortcuts

3.4 Essential R Tips and Tricks

3.4.1 R is Case Specific and Space Neutral

Everything in R is case specific. In R, upper case letters are considered completely different than their lower case counterparts. R will show no mercy if you mistype something in a difference case. However, R is pretty flexible about spaces and will usually accept multiple (or none) spaces where the default expectation is one space.

3.4.2 Comments and Line Continuation

You can add your comments/remarks in your R code by inserting the # character. Any text after the # character will be ignored by R during execution.

It is essential that sufficient comments are added to your R code so that the code remains understandable and if you share your code with somebody else then they have a way of figuring out what you are trying to do, why and how.

R is smart enough to figure when the command is not finished and allows you to continue typing on the next line. This usually works well when there are unfinished parentheses or operators at the end. R indicates that it is expecting further inputs by showing the character `+` at the beginning of the next line. See figure for an example.

```
# %![Line Continuation for Incomplete Commands](images/incomplete.png)
```

3.4.3 Naming Conventions

Object names in R can consist of any combination of alphanumeric characters (`a-z`, `A-Z`}, `0-9`) as well as underscore `_` and `.` characters. Object names should not include special characters such as `?` or space. Object names must start with a letter or a period. Special care should be taken to not include one of the reserved keywords (like `function` or `help`) as object names.

3.5 RMarkdown (RMD) Files

We will use extensive use of RMarkdown files in the class. Markdown is a simple formatting syntax for authoring documents. RMarkdown allows you to create Markdown documents with embedded R code in them. RMarkdown documents can then be rendered in variety of output formats such as HTML, PDF, and MS Word documents.

RStudio streamlines the process of creating an RMarkdown file. Getting started is easy - in RStudio, go File -> New File -> R Markdown (or Alt-F-M). This will start a new RMD file for you. It needs some startup information to proceed first, as shown in the image below:

```
# %![Figure: RMD Startup Dialog](RMDStart.png)
```

Title and Author are self explanatory. As the dialog suggests, you can choose HTML, PDF or Word as default output formats. PDF and Word require you to have a local installed version of TeX or MS-Word. For our class, I recommend that you use HTML as the default output - it is device and platform

independent - and capable of handling variety of rich media. We will see later how to include dynamic content like animation and user interaction in these HTML files as well.

Once you say OK, you will be taken to the default start page. This page has some content already to help you started. You should be able to see the preamble right on top that will generate your document headers and also control various options:

```
title: "Untitled"
author: "Dr. Sanjeev Kumar"
date: "May 5, 2018"
output: html_document
```

You can make changes by either hand-editing the information above or by clicking the Edit Option icon (which looks like a gear) and choosign your options. For example, the preamble for one of my document looks like the following:

```
title: "How To Write RMarkdown Files"
author: "Dr. Sanjeev Kumar"
date: "May 5, 2016"
output:
  html_document:
    highlight: tango
    number_sections: yes
```

Now you can go ahead and add content to your document. When you want to create the output, click the Knit button and a document will be generated that includes both content as well as the output of any embedded R code chunks within the document. We will see later in the document how to embed R code in RMD files.

It is easy to format RMarkdown files and embed R code in them. For more details on using R Markdown see <http://rmarkdown.rstudio.com>.

3.6 Embedding R Code Blocks in RMD Files

3.6.1 Embedding Options

- echo
- eval
- cache
- label

- `fig.height/width`

3.6.2 Embedding Code Blocks of Other Languages

3.6.3 Embedding in-line R Code

3.6.4 Reproducibility of Data Analysis

A key advantage of using an RMD file is that the output contains both the R code and the output. That allows for a precise matching of what steps were performed to reach a specific output. Consider the example of a complex chart made in MS-Excel - you would be hard pressed to figure out what exactly was done to create that chart. Whether the chart has some data manipulation in it, whether the chart inadvertently includes a mistake - there is no transparency - and hence you are likely to not have much trust in the chart. In the case of RMD files on the other hand, the entire process of creating the chart is right there in the form of R commands. Anybody else with the same data and the same commands will reach exactly the same output. Your analysis can be reproduced by anybody else - improving the reliability and trustworthiness of your analysis.

We discussed “veracity” as one of the four Vs of Big Data. Doing data analytics in a way that is reproducible (e.g through an RMD file) goes a long way towards establishing confidence in the results of the analysis.

3.7 Key Features of R/RStudio Ecosystem

3.7.1 Open Source Software

Open Source Software is a broad term for software that is developed by a community of volunteers and is available free of charge for you and I to use. R is open source with a strong and dedicated community that takes the stewardship of R very very seriously.

As R is open source, the underlying computer code (the source code) is available for anyone to view and modify. All these extra eyes looking at the code usually means a more reliable, error free and malware resistant software.

Open source nature of R has both positive and negative impacts. On the positive side, the community helps extend R rapidly by building new packages with new capabilities very quickly. On the negative side, the R platform is built

by experts for experts - the focus is on capability and precision - not on ease of use. This makes R a swiss army knife of that is capable of doing almost everything possible in data domain but a knife that can be difficult to learn how to use.

3.7.2 Extensability of R

R is designed to be easily extended. You, me and everyone else can write new functionalities for R and have it available for everyone in the world to download and use. The R community has been writing extensions for decades now and the result is a library of thousands of packages. If there is anything data analysis related that you want to do, chances are that somebody else wanted to do it too, has already built a package for it that you can download and use.

Easy extensability of R is one of the main reasons why R remained the most capable, cutting edge data analytical platform for decades now. R community has written packages to allow you write SQL commands, C++ commands, build geo-maps, build predictive models - all within the R platform. New developments tend to get incorporated in R much faster than in competing statistical software such as Stata, SAS or SPSS.

3.7.3 Usability of R

R is not graphical. Much of R's capabilities are accessed through typing R commands rather than clicking buttons or accessing menus. R commands allow a level of precision and flexibility that is typically now achievable through a graphical interface. Doing complex tasks successfully in R involves writing a sequence of R commands that are closer to coding in a programming language than using a point and click software.

While R can be considered to have a steep learning curve, it should be noted that once you get through the learning curve and get used to R's syntax, you can slice through complex data analysis with much greater speed and capability than competing alternatives. In my opinion, the learning curves for R and competing graphical alternatives look like the figure below. R starts slower and has a difficult jumpstart but if you stay with it for a little bit then R handily outperforms the graphical alternatives.

```
curve(x^2,from = 0,to = 3, xlab = "Time", ylab="Expertise",
      col = "red", lwd=2, xaxt = 'n', yaxt = 'n')
abline(a = 2,b = 1, col = "blue", lwd = 2)
```




FIGURE 3.4: Learning Curves for R (red) and alternatives (blue)

3.7.4 R is NOT Verbose

In fact R is very very concise. It will not give you an output until you ask for it. This is very different than most statistical software that tend to overwhelm the user with every possible output possible. R allows you the flexibility to specifically request the level of detail you need. By default, R will provide only the minimal amount of output.

3.7.5 Object Oriented Programming

3.7.6 R Community's Sense of Humor

R community is a hard working one - they are continuously building new and latest packages - making the R ecosystem even better. R community is also a fun one - a clever one at the very least. For example - look at the nickname for the current version:

```
version$nickname
```

```
## [1] "Kite-Eating Tree"
```

The R community is very helpful and will answer your questions if you post on relevant forums like StackExchange.com. However, they don't suffer fools and expect people who ask questions to have done their homework. Over the

years, the R community has developed an extensive knowledgebase that is only a google search away for you. As you go further in your exploration of R, attempt to learn from the community and hopefully one day be a bonafide member of the community.

3.8 Common Issues, Pitfalls, Bugs, Obstacles and Their Solutions

This section is my master list of common issues that students encounter when they start working with R. You may want to not read this in full right now - treat this as a reference section. When you run into an issue, come by here and see if I have a solution for you.

3.8.1 I have an older R installtion. How do I upgrade it?

You might have an old, not updated version of R in your machine. You can check the current version of your R by running the `version` command. If you have an older version and if you are in a Windows machine, then upgrading your R installation is easy with the `installr` package. The `installr` package has the handy `updateR` command that will automatically update your R installation to the latest one. A sample code to accomplish this is below:

```
#Sample code, not being evaluated here  
install.packages("installr") # install the installr package  
library(installr)  
updateR() # updating your R version  
version # command for checking your current R version
```

If you happen to be in a Mac or Linux machine, then you would want to download the current version of R and install. This process is similar to installing a fresh new installation of R as discussion in a previous section.

3.8.2 Help - I am not being able to install any packages!

If you are unable to install R packages, there are a few things we can try.² First thing to check would be your anti-virus software. Several anti-virus software

²I am not going to insult your intelligence by asking you to make sure that you have a live Internet connection!

balk at allowing R to perform tasks that may be interpreted as a security risk (for example - downloading and installing packages from the Internet). You can try to work around this by disabling/uninstalling your anti-virus or adding R to the exceptions list in the anti-virus.

You may also want to check whether you have administrator access in the machine you are working on. Several essential tasks in R require the user to have administrator access.

Lastly, R may be attempting to install packages into a directory that it does not have write permissions on. This can sometimes happen if you are on a computer on a network with mapped drives. You can work around this problem by asking R to install your packages in a safe, editable directory where you have write access. R installs all packages in a library directory - you can find out the current location of the library directory with the `.libPaths` command. You can use the `.libPaths` command to then add a directory location as your personal library location so that R can then install packages in that directory. A sample code to accomplish this is shown below:³

```
.libPaths() #Shows the current library path

## [1] "/Library/Frameworks/R.framework/Versions/3.4/Resources/library"

# You can add your own library folder to the path
#libPaths( c( .libPaths(), "C:/Whatever/Whatever..." ) )
```

3.8.3 Why is the edit command not working for me?

One of the first R commands that gives trouble is the `edit` command. As the name suggests, `edit` is used for editing data - typically it shows the data to be edited in a table format with editable cells. This specific command (and its cousin, the `fix` command) seems to be very sensitive to the keyboard layouts, especially on Apple machines. This problem is hardware oriented and may prove difficult to solve. I usually recommend to ignore these two commands and take care of your data editing needs using alternative methods. If you were using `edit` for only viewing the data, then `View` is a good alternative for that.

³Last command commented as C:/Whatever does not exist obviously - so would have resulted in an error - hence it is commented out.



A

Syllabi

This appendix will have the syllabus for the relevant Ross courses.

A.1 TO404 Big Data Manipulation and Visualization

A.2 TO414 Advanced Analytics

A.3 TO628 Advanced Big Data Analytics



Bibliography

Foundation, F. S. (2007). The gnu general public license v3.0 - gnu project - free software foundation. <https://www.gnu.org/licenses/gpl-3.0.en.html>. (Accessed on 09/23/2016).

Xie, Y. (2015). *Dynamic Documents with R and knitr*. Chapman and Hall/CRC, Boca Raton, Florida, 2nd edition. ISBN 978-1498716963.

Xie, Y. (2017). *bookdown: Authoring Books and Technical Documents with R Markdown*. R package version 0.5.



Index of R Commands, Packages and Key Concepts

bookdown, [xiii](#)

FOO, [3](#)

knitr, [xiii](#)